



Data-Intensive Distributed Computing

CS 451/651 431/631 (Winter 2018)

Part 8: Analyzing Graphs, Redux (1/2)
March 20, 2018

Jimmy Lin
David R. Cheriton School of Computer Science
University of Waterloo

These slides are available at <http://lintool.github.io/bigdata-2018w/>



This work is licensed under a Creative Commons Attribution-Noncommercial-Share Alike 3.0 United States
See <http://creativecommons.org/licenses/by-nc-sa/3.0/us/> for details

Graph Algorithms, again? (srsly?)

What makes graphs hard?

X

Fun!

Irregular structure

Fun with data structures!

Irregular data access patterns

Fun with architectures!

Iterations

Fun with optimizations!

Characteristics of Graph Algorithms

Parallel graph traversals


Local computations

Message passing along graph edges



Iterations

Visualizing Parallel BFS



PageRank: Defined


Given page x with inlinks t_1, \dots, t_n , where

$C(t)$ is the out-degree of t


α is probability of random jump

N is the total number of nodes in the graph

$$PR(x) = \alpha \left(\frac{1}{N} \right) + (1 - \alpha) \sum_{i=1}^n \frac{PR(t_i)}{C(t_i)}$$



PageRank in MapReduce



PageRank vs. BFS

PageRank

BFS

Map

PR/N

$d + I$

Reduce

sum

min


Characteristics of Graph Algorithms

Parallel graph traversals


Local computations

Message passing along graph edges


Iterations



BFS



PageRank



MapReduce Sucks


Hadoop task startup time


Stragglers


Needless graph shuffling


Checkpointing at each iteration


Let's Spark!















MapReduce vs. Spark




Characteristics of Graph Algorithms

Parallel graph traversals

Local computations

Message passing along graph edges

Iterations



Even faster?

Big Data Processing in a Nutshell

Let's be smarter about this!

Partition

Replicate

Reduce cross-partition communication


Simple Partitioning Techniques

Hash partitioning

Range partitioning on some underlying linearization


Web pages: lexicographic sort of domain-reversed URLs

How much difference does it make?




PageRank over webgraph
(40m vertices, 1.4b edges)

How much difference does it make?




PageRank over webgraph
(40m vertices, 1.4b edges)

How much difference does it make?



PageRank over webgraph
(40m vertices, 1.4b edges)

How much difference does it make?



PageRank over webgraph
(40m vertices, 1.4b edges)

Schimmy Design Pattern

Basic implementation contains two dataflows:


Messages (actual computations)


Graph structure (“bookkeeping”)

Schimmy: separate the two dataflows, shuffle only the messages


Basic idea: merge join between graph structure and messages

both relations sorted by join key
consistently partitioned and sorted by join key






How much difference does it make?



PageRank over webgraph
(40m vertices, 1.4b edges)

How much difference does it make?



PageRank over webgraph
(40m vertices, 1.4b edges)

Simple Partitioning Techniques


Hash partitioning

Range partitioning on some underlying linearization

Web pages: lexicographic sort of domain-reversed URLs

Social networks: sort by demographic characteristics

Country Structure in Facebook



Analysis of 721 million active users (May 2011)

54 countries w/ >1m active users, >50% penetration

Simple Partitioning Techniques

Hash partitioning


Range partitioning on some underlying linearization

Web pages: lexicographic sort of domain-reversed URLs


Social networks: sort by demographic characteristics

Geo data: space-filling curves





Aside: Partitioning Geo-data




Geo-data = regular graph



Space-filling curves: Z-Order Curves



Space-filling curves: Hilbert Curves



Simple Partitioning Techniques

Hash partitioning

Range partitioning on some underlying linearization

Web pages: lexicographic sort of domain-reversed URLs

Social networks: sort by demographic characteristics

Geo data: space-filling curves


But what about graphs in general?





General-Purpose Graph Partitioning

Graph coarsening
Recursive bisection

General-Purpose Graph Partitioning



Graph Coarsening



Karypis and Kumar. (1998) A Fast and High Quality Multilevel Scheme for Partitioning Irregular Graphs.

Chicken-and-Egg

To coarsen the graph you need to identify dense local regions

To identify dense local regions quickly you need traverse local edges

But to traverse local edges efficiently you need the local structure!

To efficiently partition the graph, you need to already know what the partitions are!


Industry solution?

Big Data Processing in a Nutshell


Partition

Replicate


Reduce cross-partition communication



Partition



Partition



What's the fundamental issue?


Characteristics of Graph Algorithms

Parallel graph traversals


Local computations

Message passing along graph edges


Iterations



Partition




State-of-the-Art Distributed Graph Algorithms



Graph Processing Frameworks





Pregel: Computational Model

Based on Bulk Synchronous Parallel (BSP)

Computational units encoded in a directed graph

Computation proceeds in a series of supersteps

Message passing architecture

Each vertex, at each superstep:

Receives messages directed at it from previous superstep

Executes a user-defined function (modifying state)


Emits messages to other vertices (for the next superstep)

Termination:

A vertex can choose to deactivate itself

Is “woken up” if new messages received

Computation halts when all vertices are inactive



Pregel: Implementation

Master-Worker architecture

Vertices are hash partitioned (by default) and assigned to workers
Everything happens in memory

Processing cycle:

Master tells all workers to advance a single superstep
Worker delivers messages from previous superstep, executing vertex computation
Messages sent asynchronously (in batches)
Worker notifies master of number of active vertices

Fault tolerance

Checkpointing
Heartbeat/revert

Pregel: SSSP

```
class ShortestPathVertex : public Vertex<int, int, int> {
    void Compute(MessageIterator* msgs) {
        int mindist = IsSource(vertex_id()) ? 0 : INF;
        for (; !msgs->Done(); msgs->Next())
            mindist = min(mindist, msgs->Value());
        if (mindist < GetValue()) {
            *MutableValue() = mindist;
            OutEdgeIterator iter = GetOutEdgeIterator();
            for (; !iter.Done(); iter.Next())
                SendMessageTo(iter.Target(),
                              mindist + iter.GetValue());
        }
        VoteToHalt();
    }
};
```


Pregel: PageRank

```
class PageRankVertex : public Vertex<double, void, double> {
public:
    virtual void Compute(MessageIterator* msgs) {
        if (superstep() >= 1) {
            double sum = 0;
            for (; !msgs->Done(); msgs->Next())
                sum += msgs->Value();
            *MutableValue() = 0.15 / NumVertices() + 0.85 * sum;
        }

        if (superstep() < 30) {
            const int64 n = GetOutEdgeIterator().size();
            SendMessageToAllNeighbors(GetValue() / n);
        } else {
            VoteToHalt();
        }
    }
};
```

Pregel: Combiners

```
class MinIntCombiner : public Combiner<int> {  
    virtual void Combine(MessageIterator* msgs) {  
  
        int mindist = INF;  
        for (; !msgs->Done(); msgs->Next())  
            mindist = min(mindist, msgs->Value());  
        Output("combined_source", mindist);  
    }  
};
```



A P A C H E
G I R A P H

Giraph Architecture

Master – Application coordinator

Synchronizes supersteps

Assigns partitions to workers before superstep begins

Workers – Computation & messaging


Handle I/O – reading and writing the graph

Computation/messaging of assigned partitions

ZooKeeper


Maintains global application state

Giraph Dataflow




Giraph Lifecycle

Vertex Lifecycle




Giraph Lifecycle




Giraph Example


```
public class MaxComputation extends BasicComputation<IntWritable, IntWritable,  
    NullWritable, IntWritable> {  
    @Override  
    public void compute(Vertex<IntWritable, IntWritable, NullWritable> vertex,  
        Iterable<IntWritable> messages) throws IOException  
{  
    boolean changed = false;  
    for (IntWritable message : messages) {  
        if (vertex.getValue().get() < message.get()) {  
            vertex.setValue(message);  
            changed = true;  
        }  
    }  
    if (getSuperstep() == 0 || changed) {  
        sendMessageToAllEdges(vertex, vertex.getValue());  
    }  
    vertex.voteToHalt();  
}
```

Execution Trace






State-of-the-Art Distributed Graph Algorithms



Graph Processing Frameworks



GraphX: Motivation



GraphX = Spark for Graphs


Integration of record-oriented and graph-oriented processing

Extends RDDs to Resilient Distributed Property Graphs

```
class Graph[VD, ED] {  
    val vertices: VertexRDD[VD]  
    val edges: EdgeRDD[ED]  
}
```

Property Graph: Example

Property Graph



Vertex Table


| Id | Property (V) |
|----|-----------------------|
| 3 | (rxin, student) |
| 7 | (jgonzal, postdoc) |
| 5 | (franklin, professor) |
| 2 | (istoica, professor) |

Edge Table


| SrcId | DstId | Property (E) |
|-------|-------|--------------|
| 3 | 7 | Collaborator |
| 5 | 3 | Advisor |
| 2 | 5 | Colleague |
| 5 | 7 | PI |

Underneath the Covers

Property Graph







Vertex Table (RDD)



Routing Table (RDD)



Edge Table (RDD)



GraphX Operators

“collection” view

```
val vertices: VertexRDD[VD]
val edges: EdgeRDD[ED]
val triplets: RDD[EdgeTriplet[VD, ED]]
```


Transform vertices and edges

```
mapVertices
mapEdges
mapTriplets
```

Join vertices with external table

Aggregate messages within local neighborhood

Pregel programs



A photograph of a traditional Japanese rock garden. In the foreground, a gravel path is raked into fine, parallel lines. Several large, dark, irregular stones are scattered across the garden. A small, shallow pond is visible in the middle ground, surrounded by more stones and some low-lying green plants. In the background, there are more stones, some small trees, and a traditional wooden building with a tiled roof. The overall atmosphere is peaceful and minimalist.

Questions?