



Data-Intensive Distributed Computing

CS 451/651 431/631 (Winter 2018)

Part 8: Analyzing Graphs, Redux (2/2)
March 22, 2018

Jimmy Lin
David R. Cheriton School of Computer Science
University of Waterloo

These slides are available at <http://lintool.github.io/bigdata-2018w/>



This work is licensed under a Creative Commons Attribution-Noncommercial-Share Alike 3.0 United States
See <http://creativecommons.org/licenses/by-nc-sa/3.0/us/> for details

Theme for Today:

How things work in the real world
(forget everything I told you...)

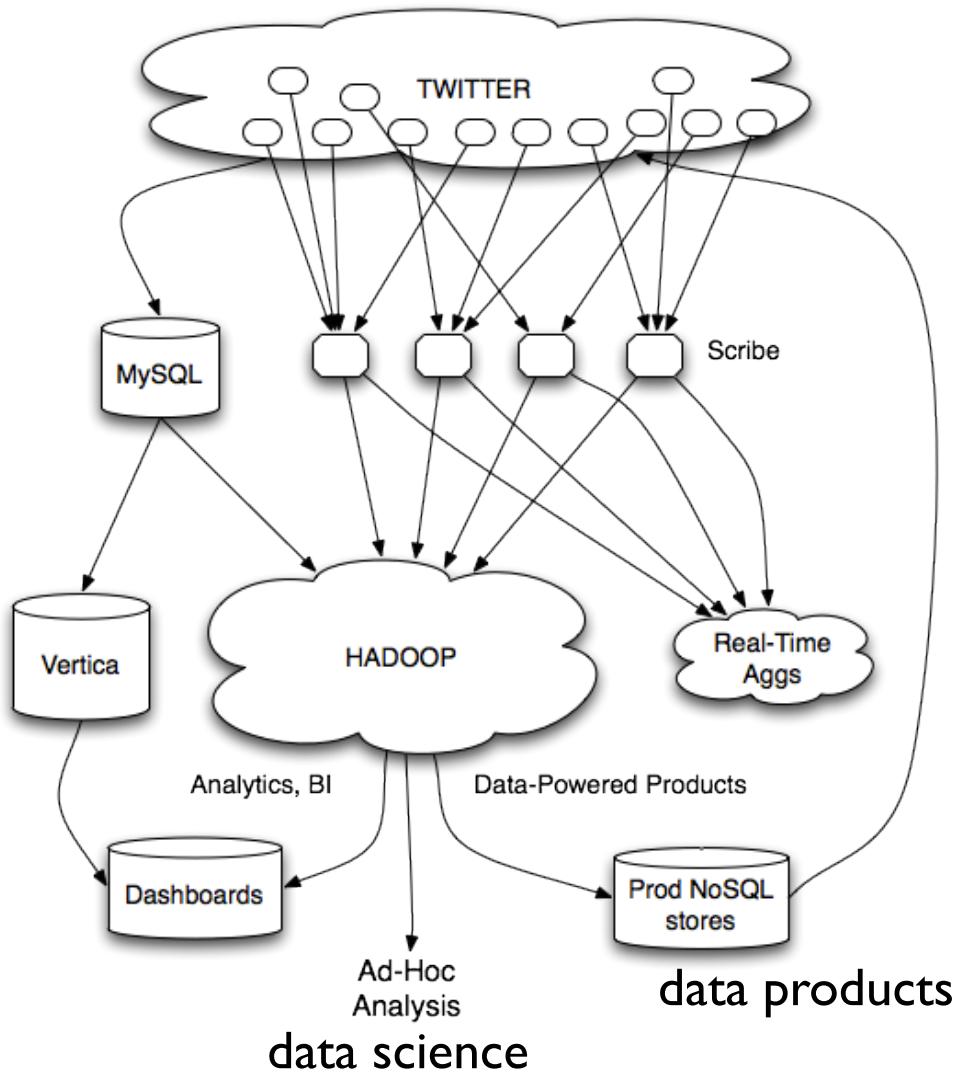


From the Ivory Tower...



... to building sh*t that works

What exactly did I do at Twitter?



I worked on...

- analytics infrastructure to support data science**
- data products to surface relevant content to users**



Tweets

Mishne et al. Fast Data in the Era of Big Data: Twitter's Real-Time Related Query Suggestion Architecture. SIGMOD 2013.



Struggling with complex data
of Data Science 2/20 to rehi

↗ Promoted by Cloudera

TWEETS FOLLOWING FOLLOWERS

1,64 Leibert et al. Automatic Management of Partitioned,
Replicated Search Services. SoCC 2011

Compose new Tweet...

Who to follow · Refresh · View all

plotly @plotlygraphs

+ Follow

↗ Promoted



Brad Anderson @boorad

Followed by Florian Leibert ...

+ Follow



Sheila Morrissey @sheilaMorr

+ Follow

Popular accounts · Find friends

Trends

#Olympics

Ukraine

#ConfessYourUnpopularOpinion

Venny

#PremioLoNuestro

I worked on...

– analytics infrastructure to support data science

– data products to surface relevant content to users

sochi



#Sochi2014

#SochiProblems

Sochi

#SochiFail

Sochi 2014 ✅ @Sochi2014

Sochi Olympics 2014 @2014Sochi

Игры Сочи 2014 ✅ @sochi2014_ru

Sochi Problems @SochiProblem

NYT Olympics @SochiNYT

Sochi Problems @SochiProblems

Search all people for sochi

Reply Retweet Favorite More

Gupta et al. WTF: The Who to Follow Service at Twitter. WWW 2013
Lin and Kolcz. Large-Scale Machine Learning at Twitter. SIGMOD 2012
@mollyhepp and @BobCusack

View summary

More

Reply Retweet Favorite More

Expand

More



circa ~2010

~150 people total

~60 Hadoop nodes

~6 people use analytics stack daily

circa ~2012

~1400 people total

10s of Ks of Hadoop nodes, multiple DCs

10s of PBs total Hadoop DW capacity

~100 TB ingest daily

dozens of teams use Hadoop daily

10s of Ks of Hadoop jobs daily

WTF

((whoontollow))

Who to follow · [refresh](#) · [view all](#)



freshbooks FreshBooks · [Follow](#)

Promoted · Followed by @zappos and others.

×



alanwarming Alan Warming · [Follow](#)

Followed by @fredwilson and others.

×



Mozzie21 Moises Henriques · [Follow](#)

can eat

×

Similar to @ryanhall3 · [view all](#)



RunnerSpace_com RunnerSpace.com · [Follow](#)

RunnerSpace.com has the latest in news and media...



chrislieto chris lieto · [Follow](#)

Chris Lieto is a top ranked World Class Triathlete, ...



runningtimes runningtimes · [Follow](#)

#numbers

(Second half of 2012)

~175 million active users

~20 billion edges


42% edges bidirectional

Avg shortest path length: 4.05

40% as many unfollows as follows daily

WTF responsible for ~1/8 of the edges

Graphs are core to Twitter



Graph-based recommendation systems
Why? Increase engagement!

A photograph of a paved path through tall, golden-brown grass. The path leads towards a range of hills in the distance under a cloudy sky.

The Journey

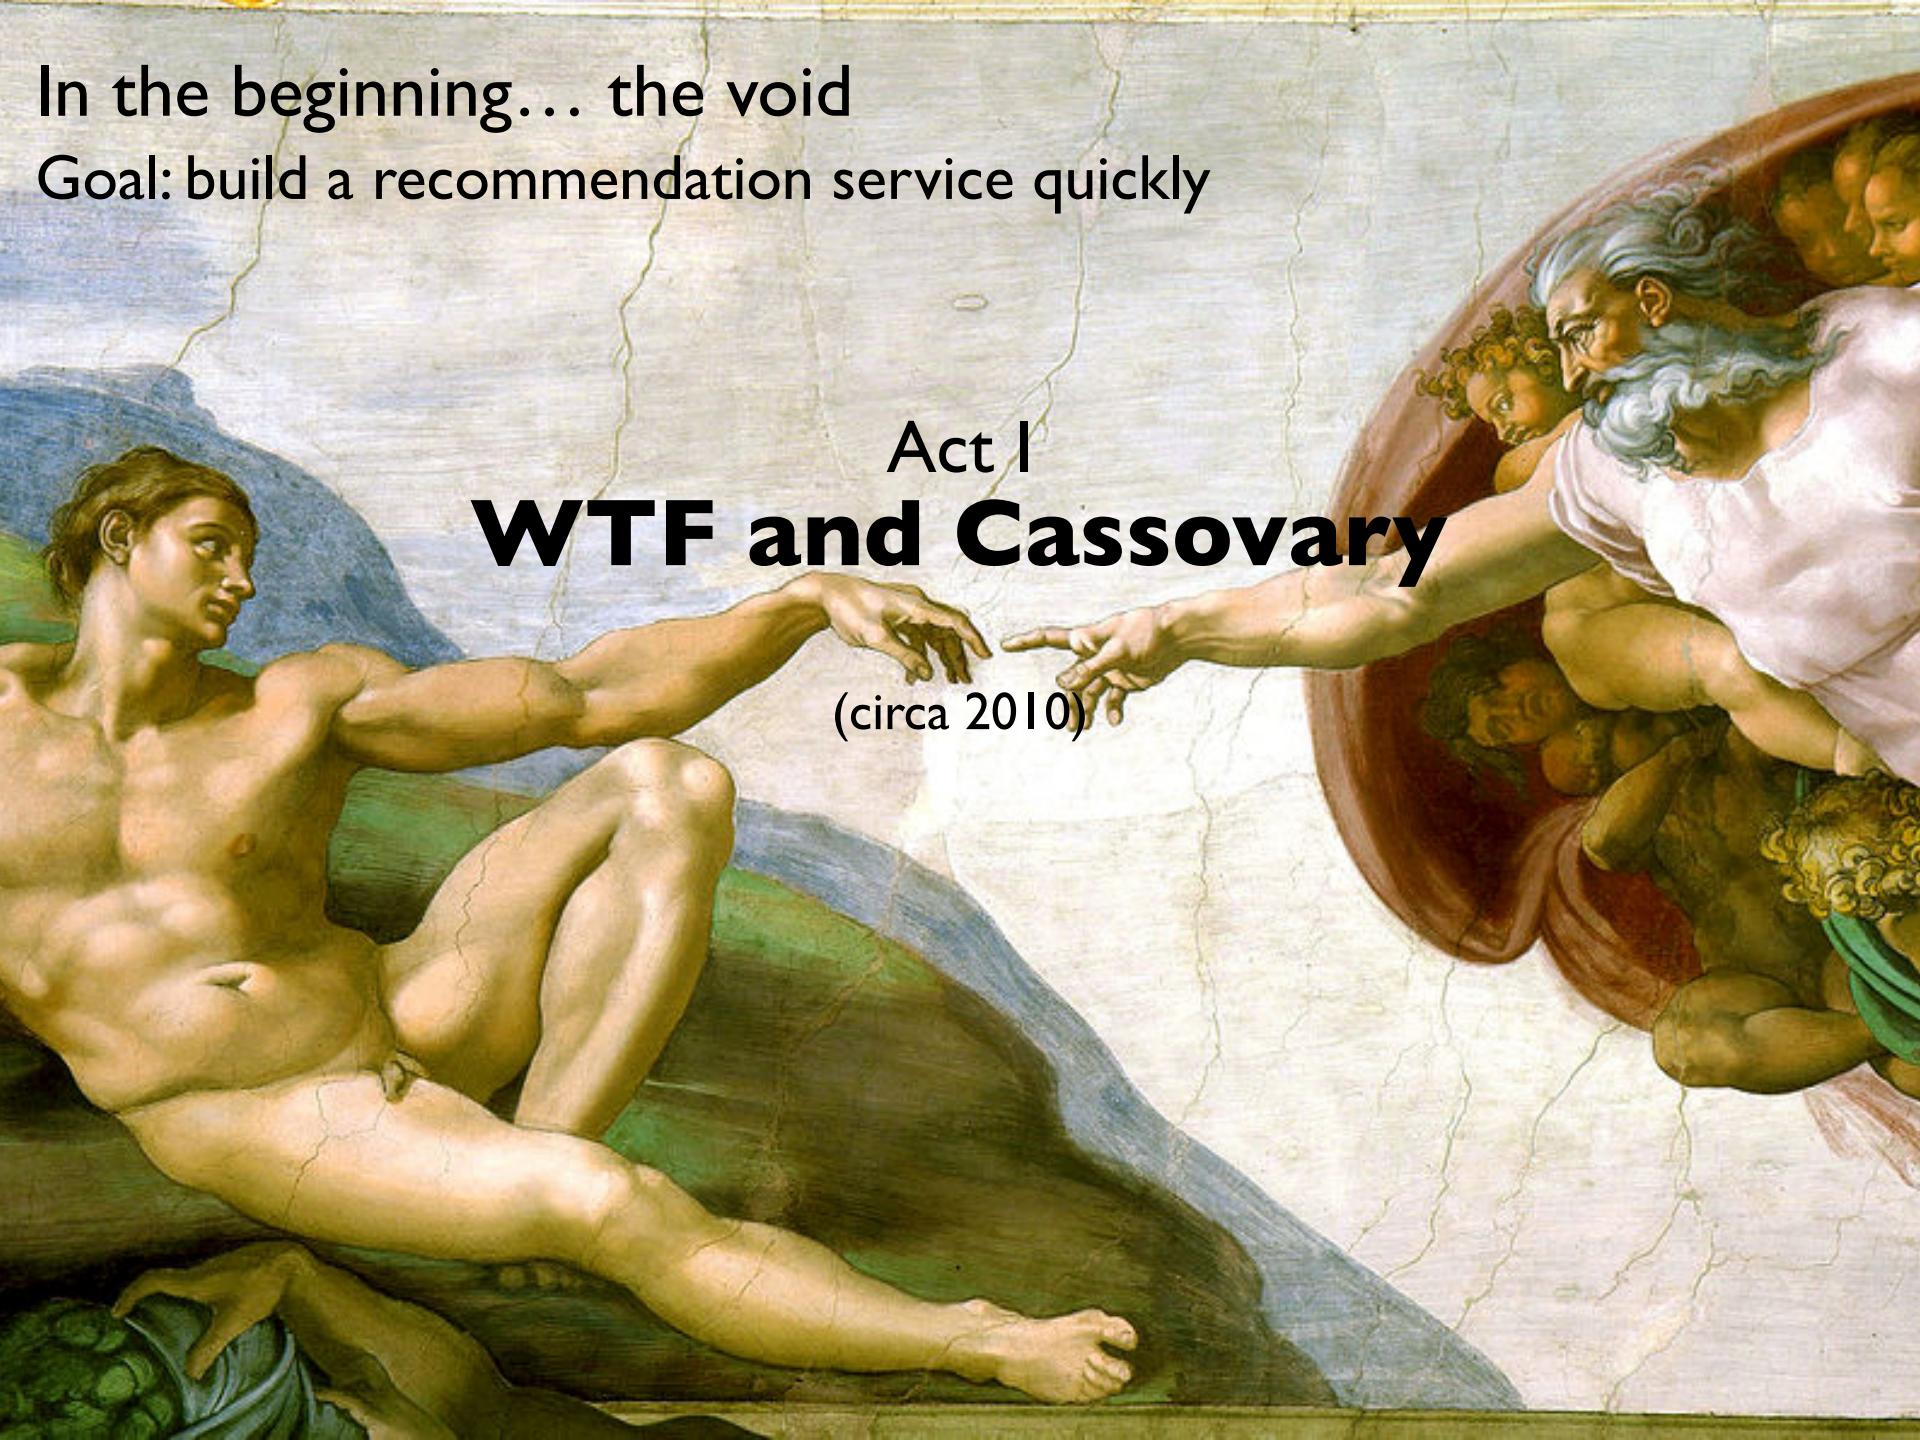
From the static follower graph for account recommendations...
... to the real-time interaction graph for content recommendations

In Four Acts...

In the beginning... the void

Act I
WTF and Cassovary

(circa 2010)



In the beginning... the void

Goal: build a recommendation service quickly

Act I

WTF and Cassovary

(circa 2010)

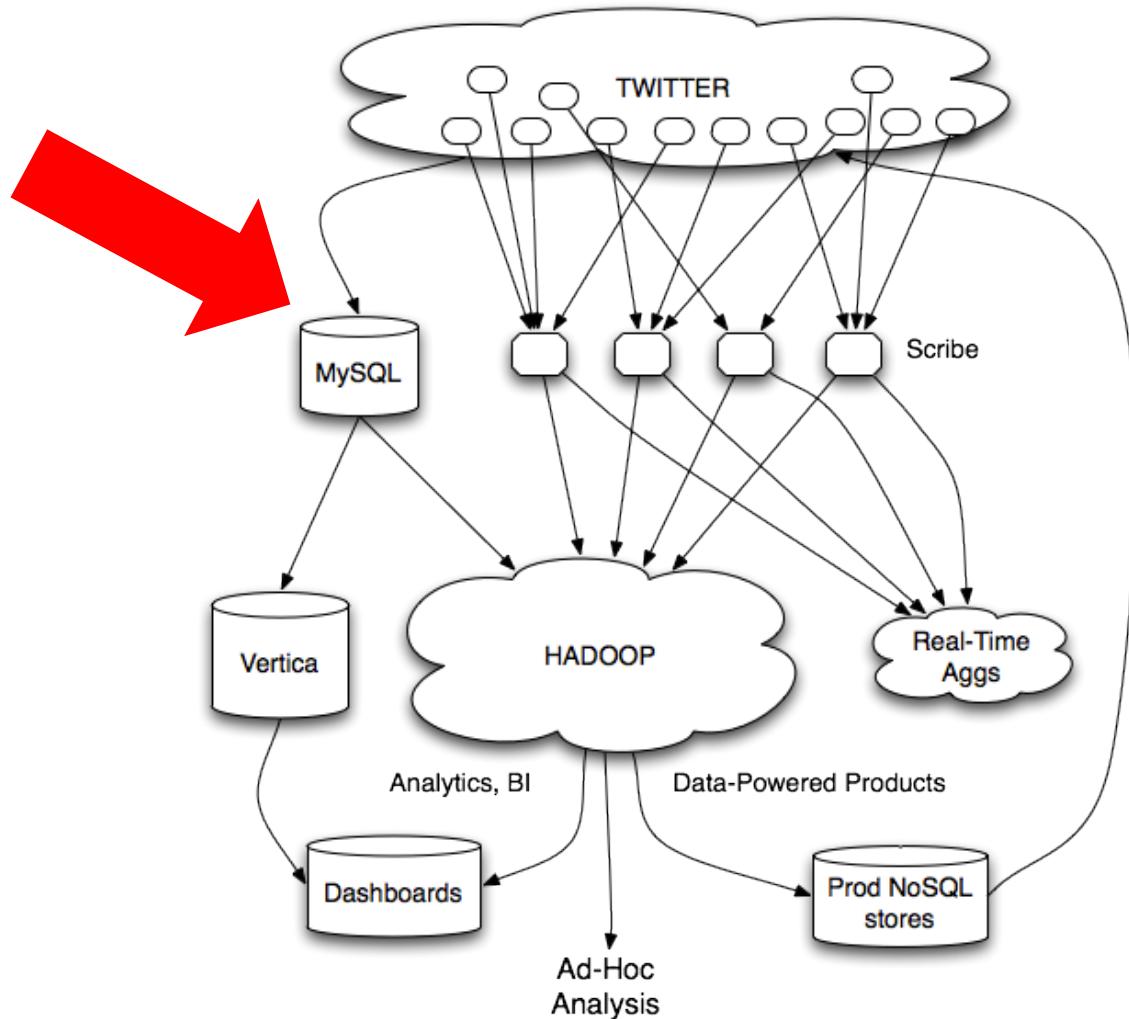


flockDB

(graph database)

Simple graph operations
Set intersection operations

Not appropriate for graph algorithms!



Okay, let's use MapReduce!
But MapReduce sucks for graphs!

What about...?

HaLoop (VLDB 2010)

Twister (MapReduce Workshop 2010)

Pregel/Giraph (SIGMOD 2010)

Graphlab (UAI 2010)

PrIter (SoCC 2011)

Datalog on Hyracks (Tech report, 2012)

Spark/GraphX (NSDI 2012, arXiv 2014)

PowerGraph (OSDI 2012)

GRACE (CIDR 2013)

Mizan (EuroSys 2013)

...

**MapReduce sucks for graph algorithms...
Let's build our own system!**

**Key design decision:
Keep entire graph in memory... on a single machine!**

Nuts!

Why?

Because we can!

Graph partitioning is hard... so don't do it
Simple architecture

Right choice at the time!

The runway argument



Suppose: 10×10^9 edges
(src, dest) pairs: ~80 GB

18×8 GB DIMMS = 144 GB

18×16 GB DIMMS = 288 GB

12×16 GB DIMMS = 192 GB

12×32 GB DIMMS = 384 GB

Cassovary

In-memory graph engine

Implemented in Scala

Compact in-memory representations

But no compression

Avoid JVM object overhead!

Open-source



PageRank

“Semi-streaming” algorithm

Keep vertex state in memory, stream over edges

Each pass = one PageRank iteration

Bottlenecked by memory bandwidth

Convergence?

Don't run from scratch... use previous values

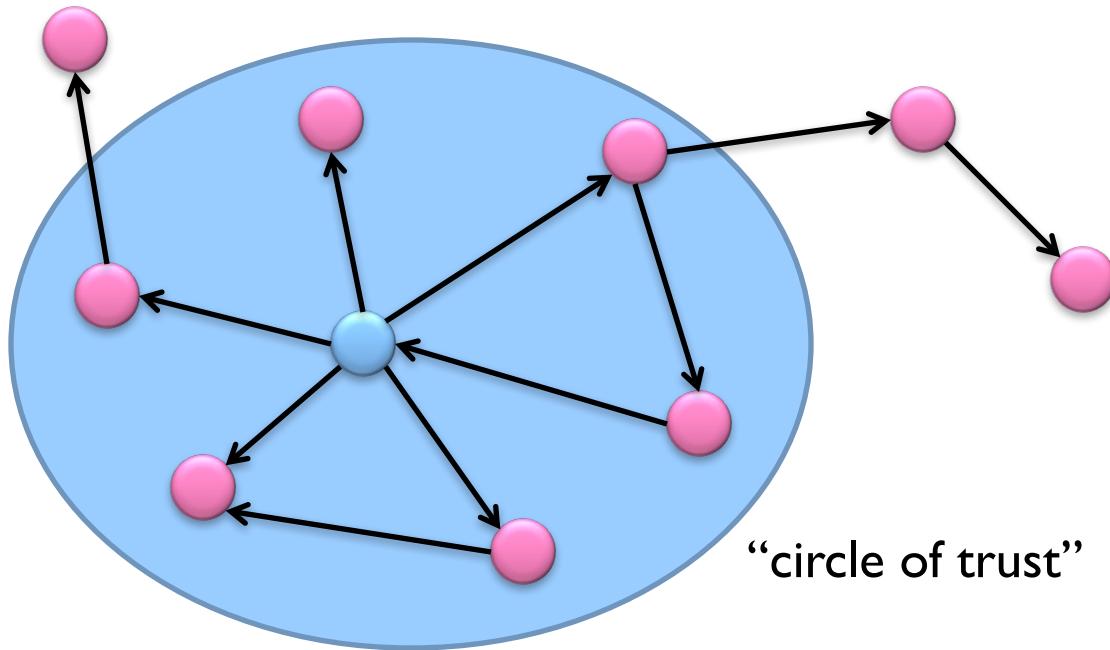
A few passes are sufficient

“Circle of Trust”

Ordered set of important neighbors for a user

Result of egocentric random walk: Personalized PageRank!

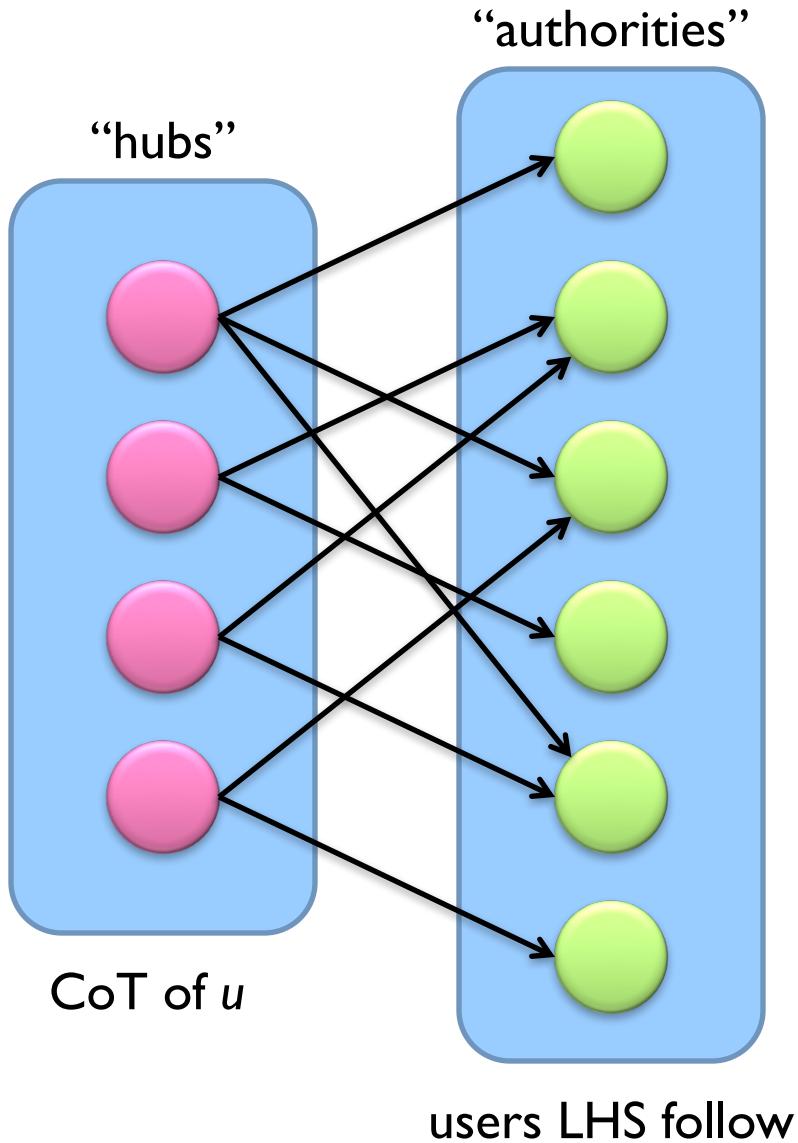
Computed online based on various input parameters



“circle of trust”


One of the features used in search

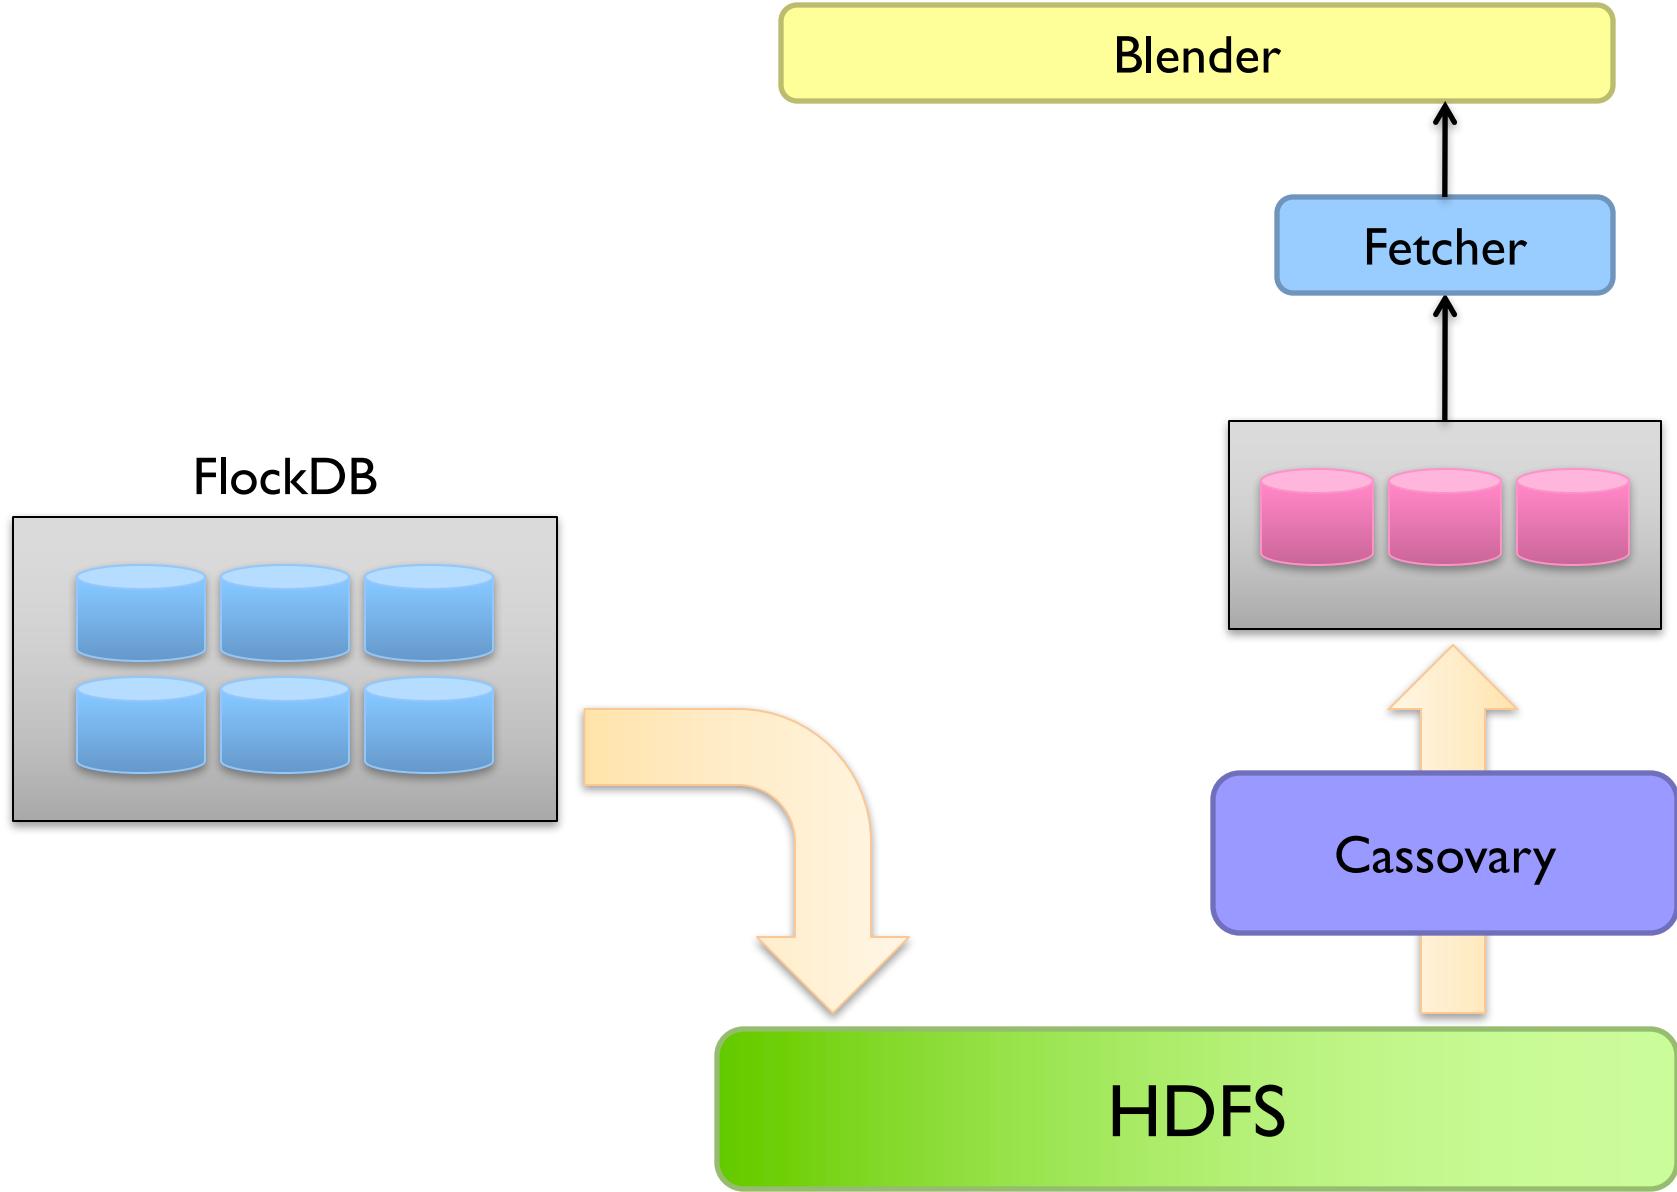
SALSA for Recommendations



hubs scores:
similarity scores to u


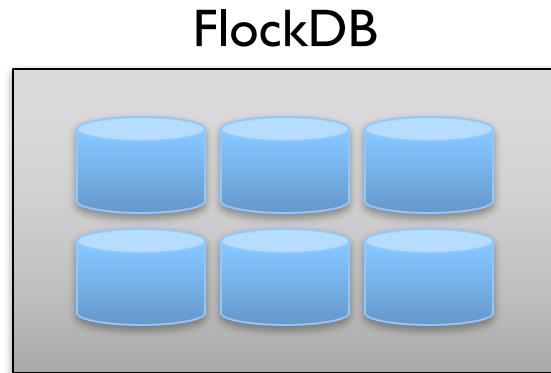
authority scores:
recommendation scores for u

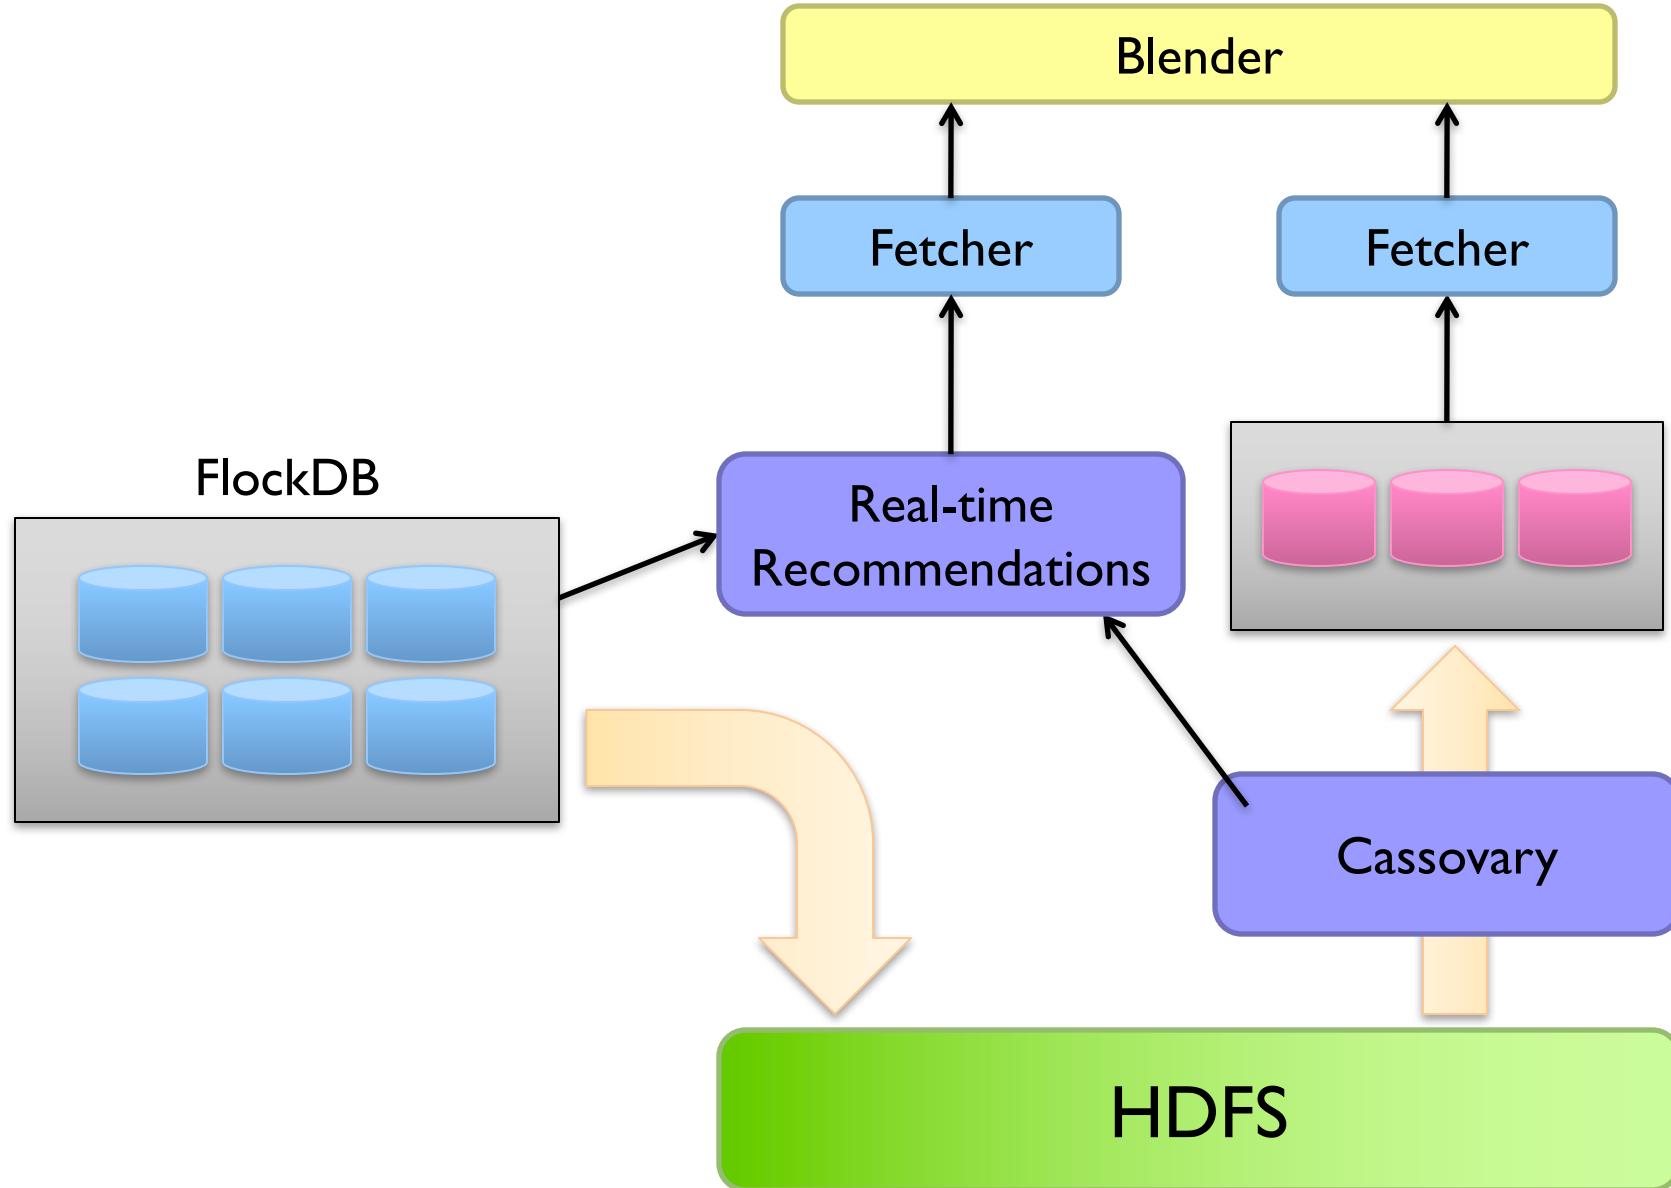




What about new users?

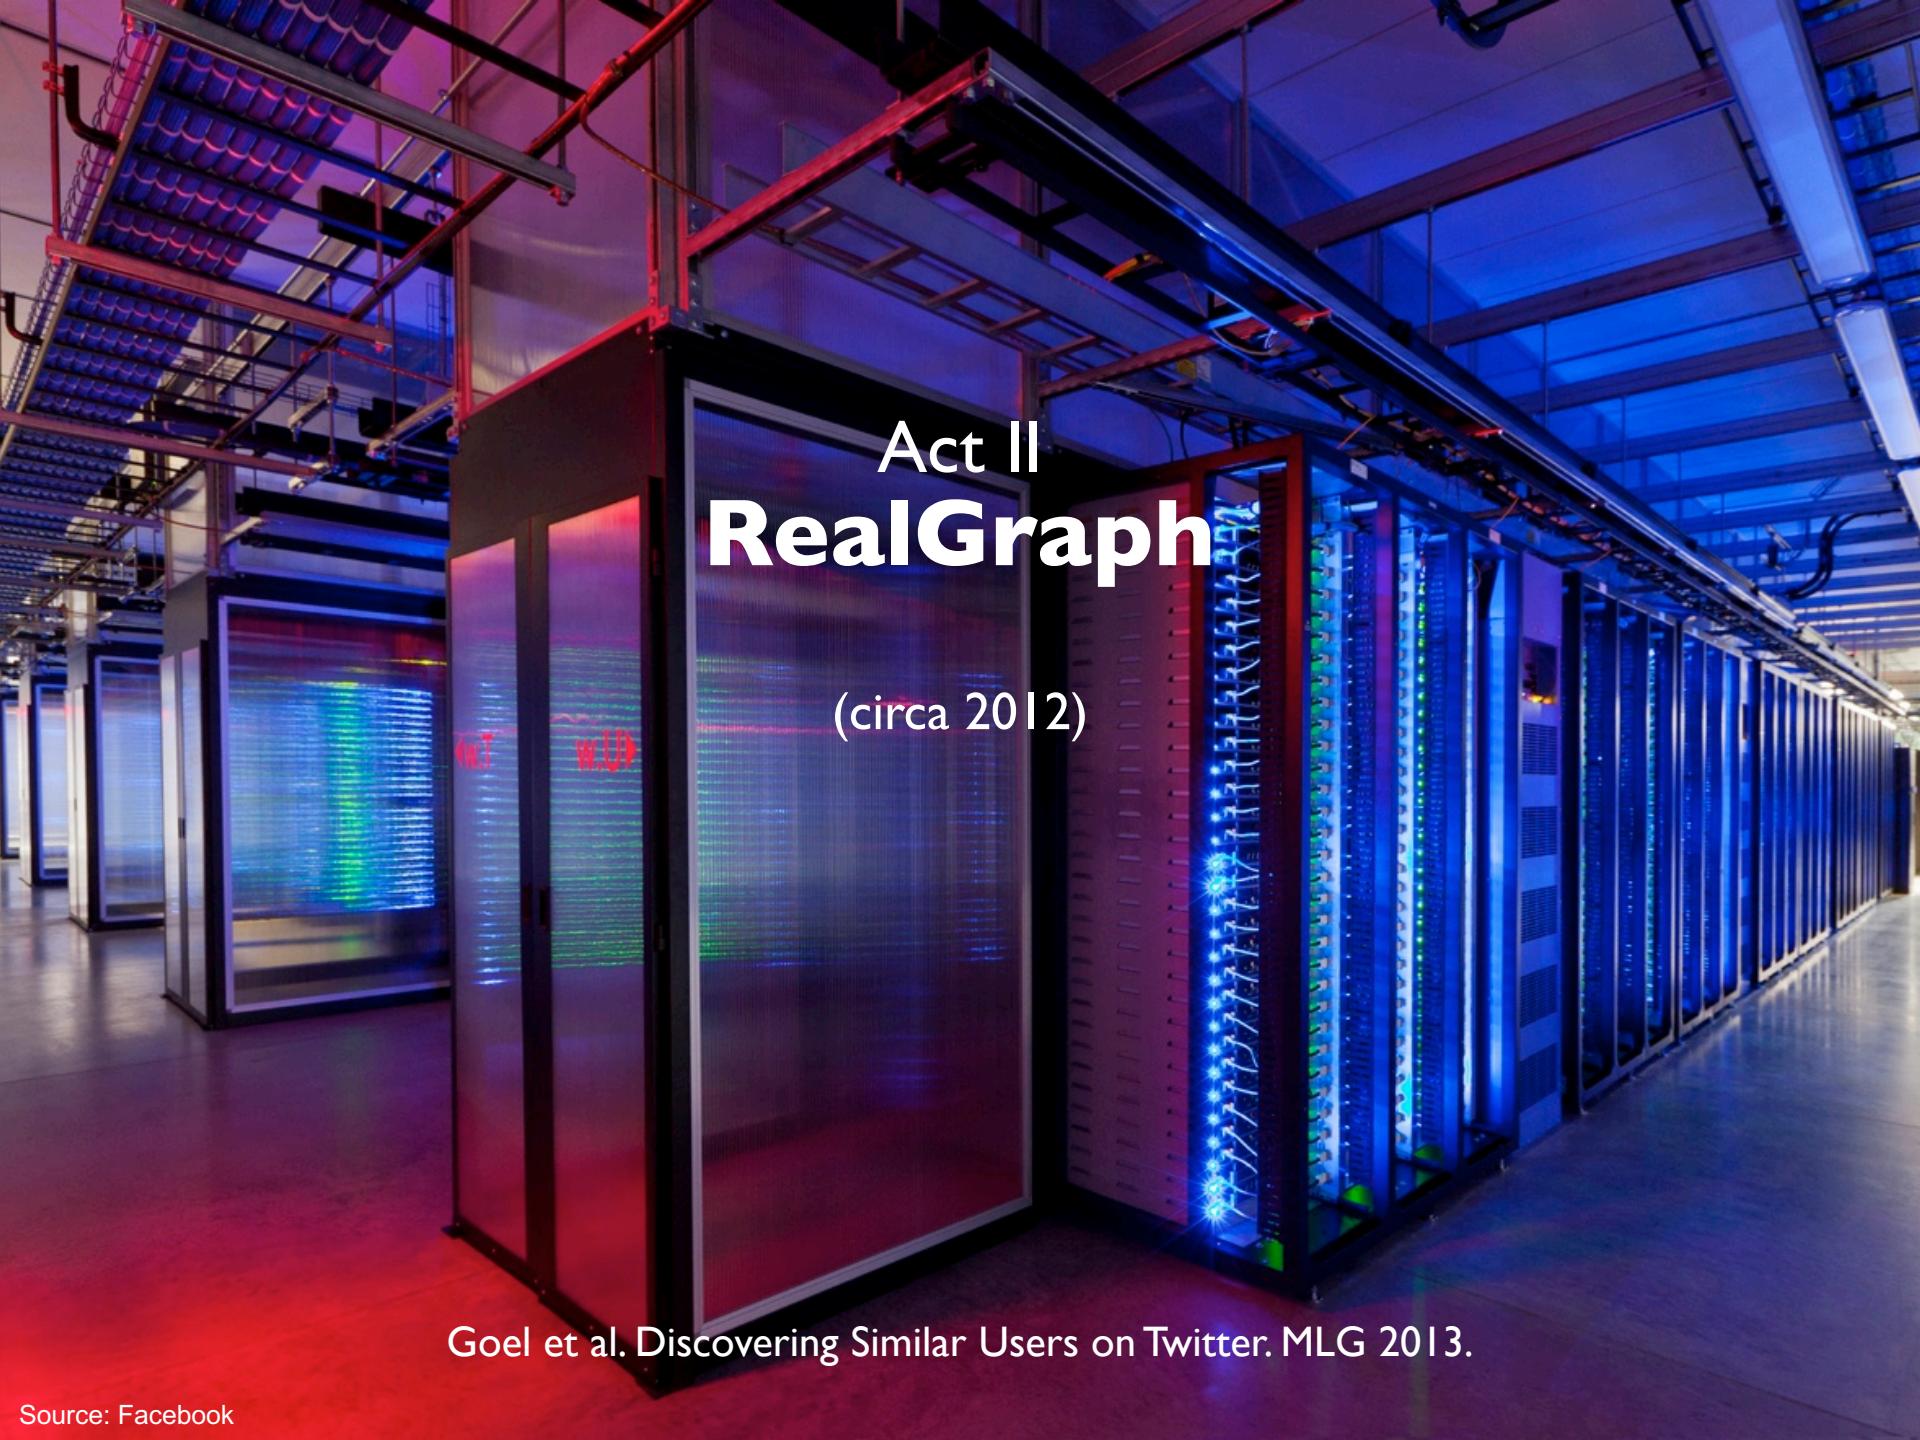
Cold start problem: they need recommendations the most!





**Spring 2010: no WTF
seriously, WTF?**

Summer 2010:WTF launched

A perspective view of a long corridor in a data center. On both sides, there are tall, dark server racks. The lighting is dramatic, with red lights on the left and blue lights on the right, creating a strong color contrast. The ceiling is white with some structural elements and pipes.

Act II RealGraph

(circa 2012)

Goel et al. Discovering Similar Users on Twitter. MLG 2013.



Another “interesting” design choice:
We migrated from Cassovary back to Hadoop!

Whaaaaaa?

Cassovary was a stopgap!

Hadoop provides:

Richer graph structure

Simplified production infrastructure

Scaling and fault-tolerance “for free”

Right choice at the time!

Wait, didn't you say MapReduce sucks?

What exactly is the issue?

Random walks on egocentric 2-hop neighborhood

Naïve approach: self-joins to materialize, then run algorithm

The shuffle is what kills you!

Graph algorithms in MapReduce

Tackle the shuffling problem!

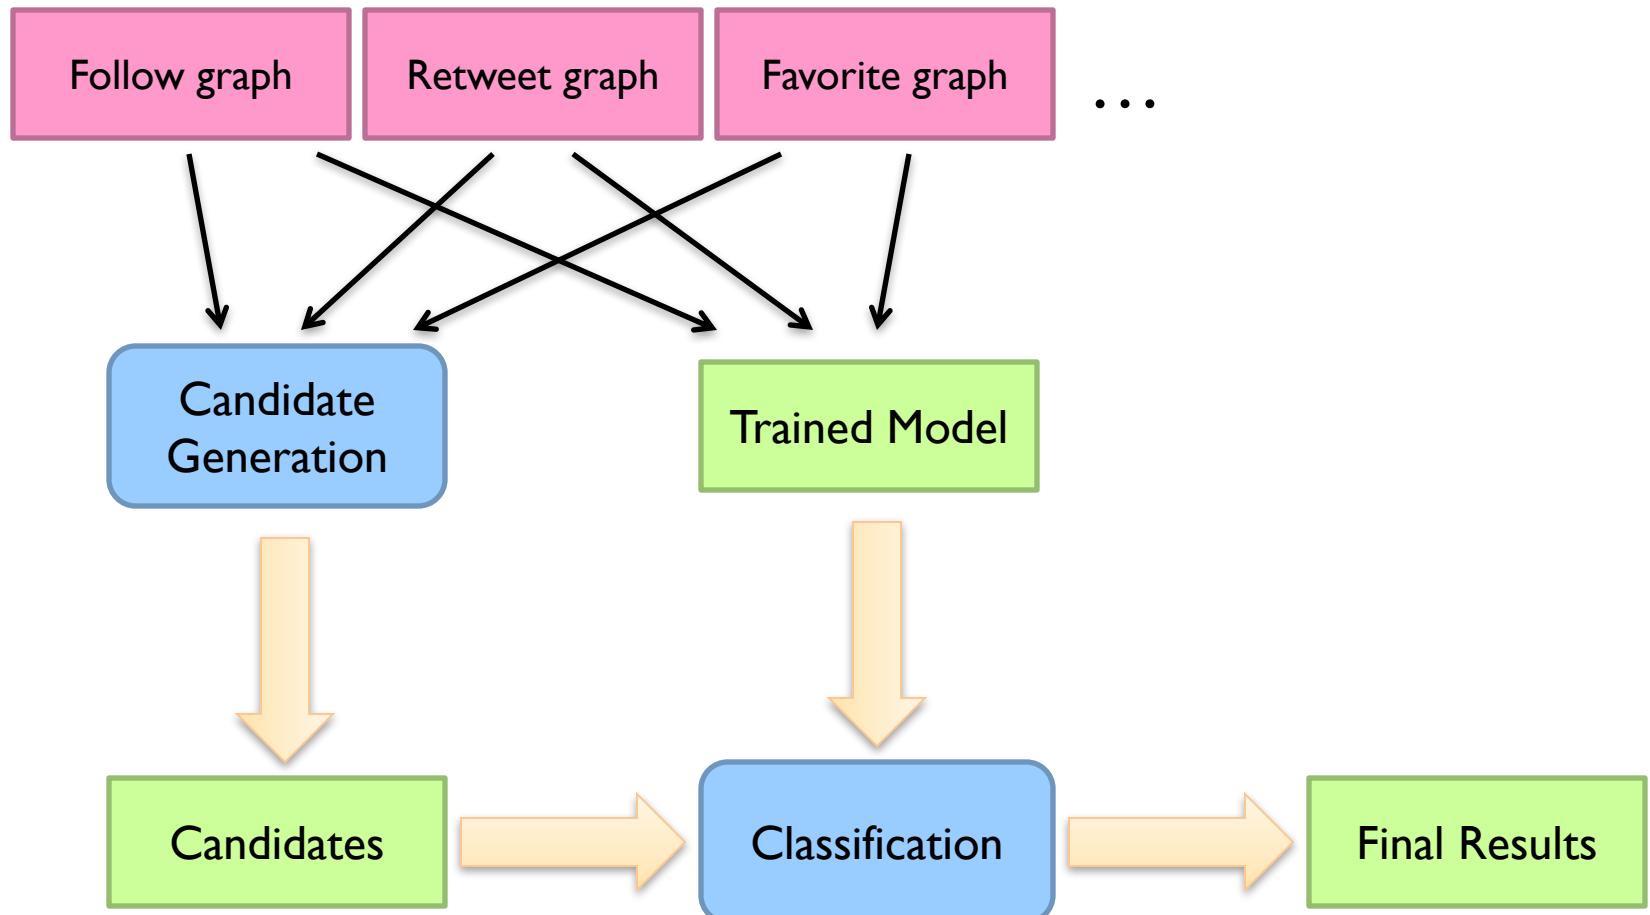
Key insights:

Batch and “stich together” partial random walks*

Clever sampling to avoid full materialization

* Sarma et al. Estimating PageRank on Graph Streams. PODS 2008
Bahmani et al. Fast Personalized PageRank on MapReduce. SIGMOD 2011.

Throw in ML while we're at it...





Act III

MagicRecs

(circa 2013)



Isn't the point of Twitter real-time?

So why is WTF still dominated by batch processing?

Observation: fresh recommendations get better engagement

Logical conclusion: generate recommendations in real time!


From batch to real-time recommendations:

Recommendations based on recent activity
“Trending in your network”

Inverts the WTF problem:

For this user, what recommendations to generate?

Given this new edge, which user to make recommendations to?



Why does this work?

A follows B's because they're interesting
B's following C's because “something's happening”
(generalizes to any activity)

Scale of the Problem

$O(10^8)$ vertices, $O(10^{10})$ edges

Designed for $O(10^4)$ events per second

Naïve solutions:

Poll each vertex periodically

Materialize everyone's two-hop neighborhood, intersect

Production solution:


Idea #1: Convert problem into adjacency list intersection

Idea #2: Partition graph to eliminate non-local intersections

Single Node Solution

Who we're recommending



“influencers”



Who we're making the recommendations to

D “dynamic” structure:
stores inverted adjacency lists
query C, return all B's that link to it


S “static” structure:
stores inverted adjacency lists
query B, return all A's that link to it



Algorithm

Who we're recommending

“influencers”



Who we're making the recommendations to

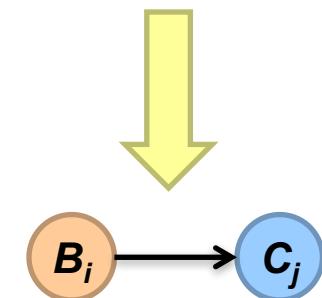
D “dynamic” structure:
stores inverted adjacency lists
query C, return all B's that link to it

1. Receive B_3 to C_2
2. Query D for C_2 , get B_1, B_2, B_3
3. For each B_1, B_2, B_3 , query S
4. Intersect lists to compute A's

S “static” structure:
stores inverted adjacency lists
query B, return all A's that link to it


Idea #1: Convert problem into adjacency list intersection

Distributed Solution



Who we're recommending


Replicate on every node



Who we're making the recommendations to

1. Fan out new edge to every node
2. Run algorithm on each partition
3. Gather results from each partition

Partition by A



Idea #2: Partition graph to eliminate non-local intersections

Production Status

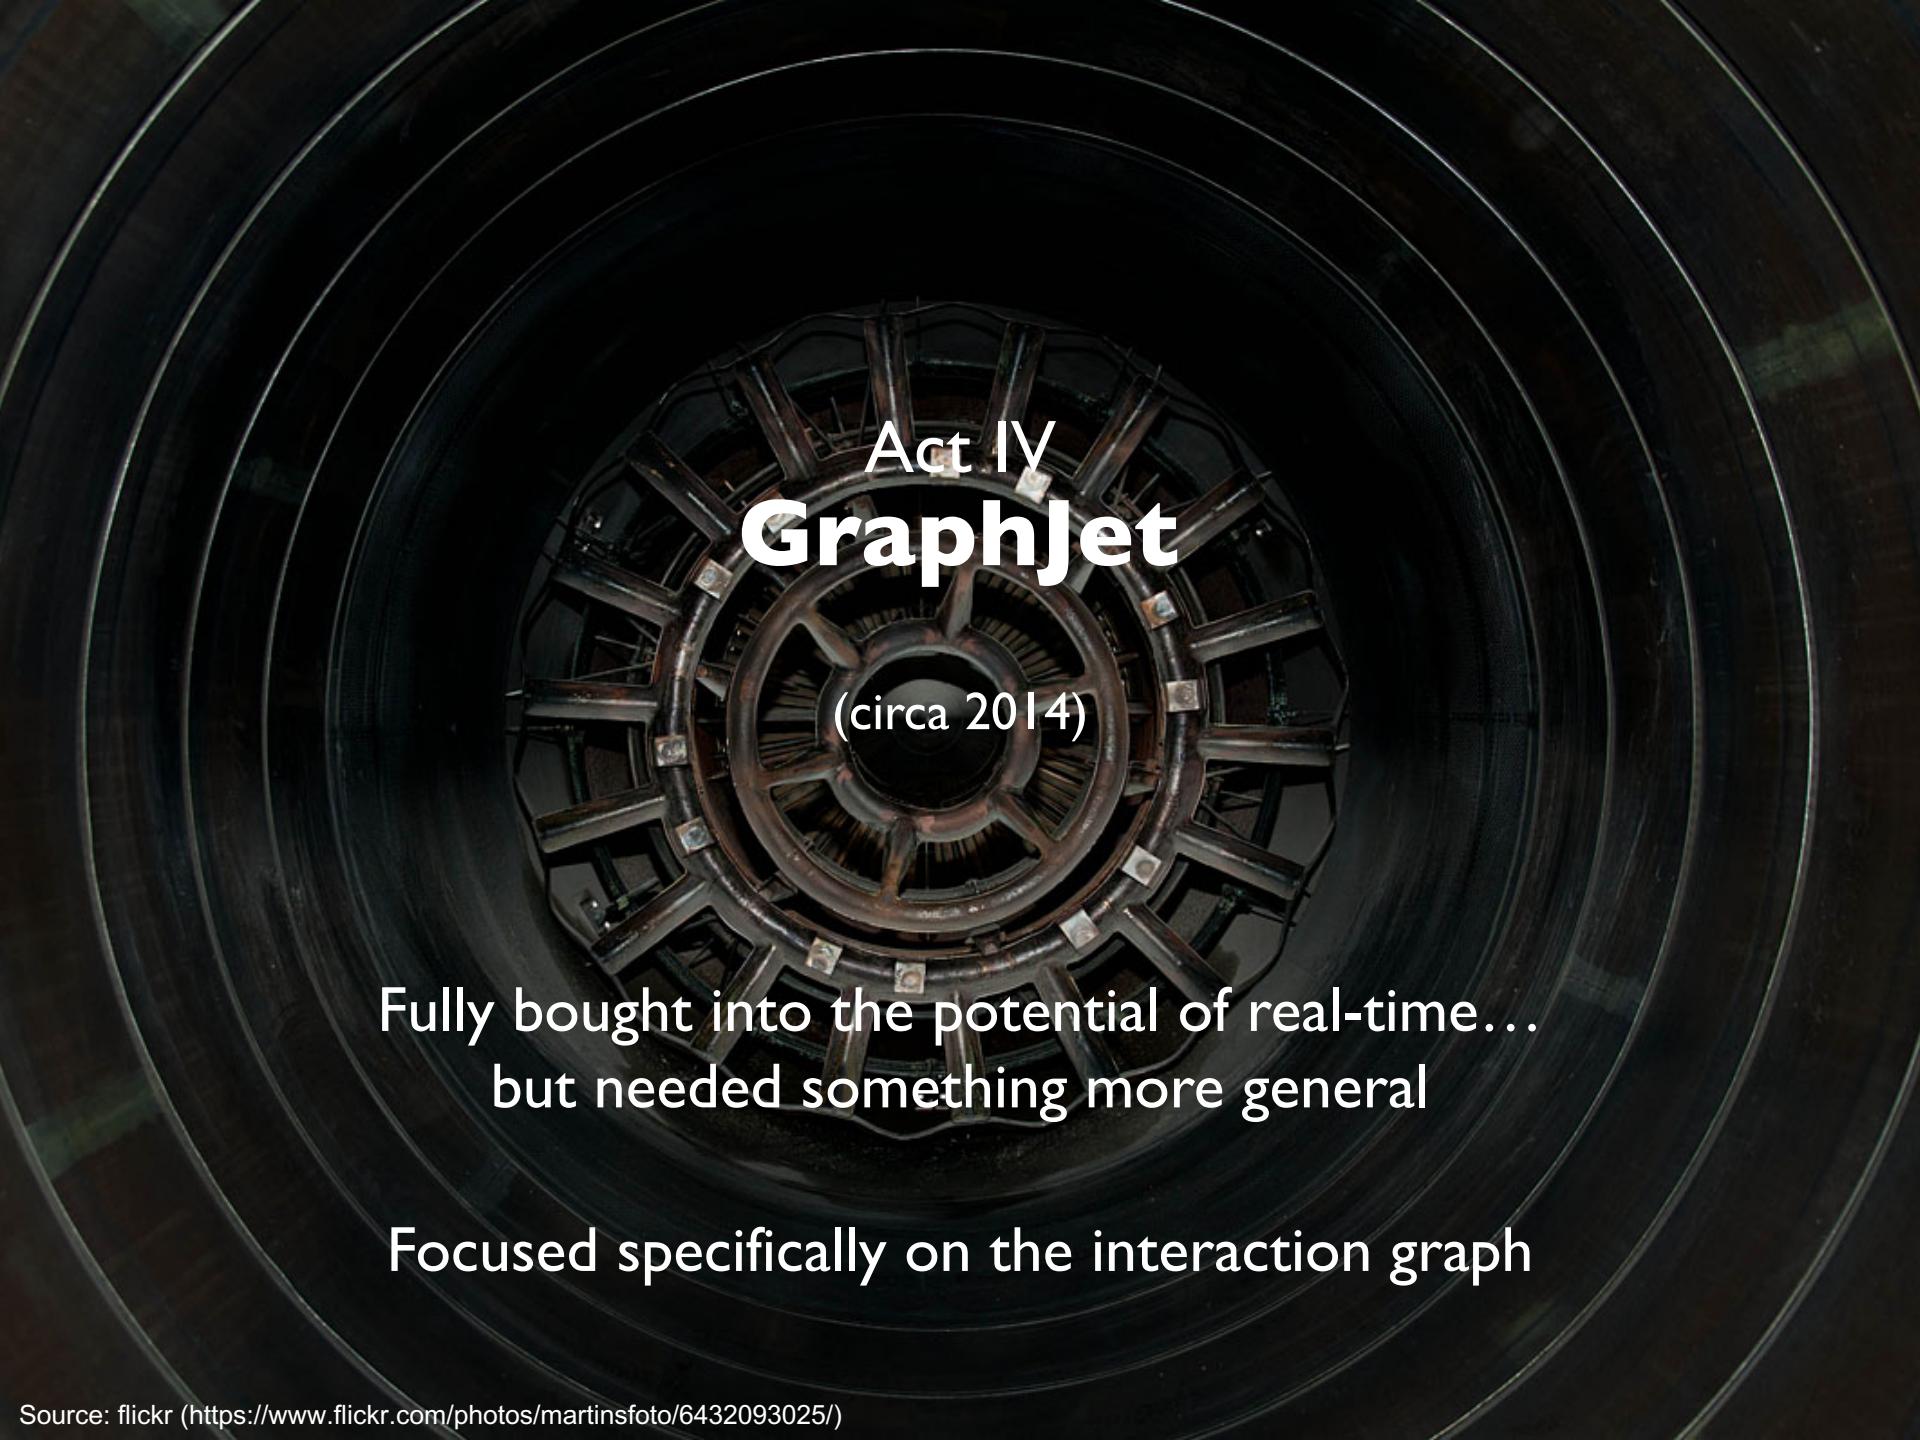
Launched September 2013

Usage Statistics (Circa 2014)

Push recommendations to Twitter mobile users
Billions of raw candidates, millions of push notifications daily

Performance

End-to-end latency (from edge creation to delivery):
median 7s, p99 15s




Act IV **GraphJet**

(circa 2014)

Fully bought into the potential of real-time...
but needed something more general

Focused specifically on the interaction graph

Data Model



Noteworthy design decisions

Make it simple, make it fast!

No partitioning

Focus on recent data, fits on a single machine

No deletes

Not meaningful w/ interaction data

No arbitrary edge metadata

Marginally better results at the cost of space – not worthwhile

Note: design supports revisiting these choices



requests

API Endpoint

Recommendation Engine

Storage Engine

getLeftVertexEdges
getLeftVertexRandomEdges
...

Index

Index Segment

Index Segment

Index Segment

Index Segment

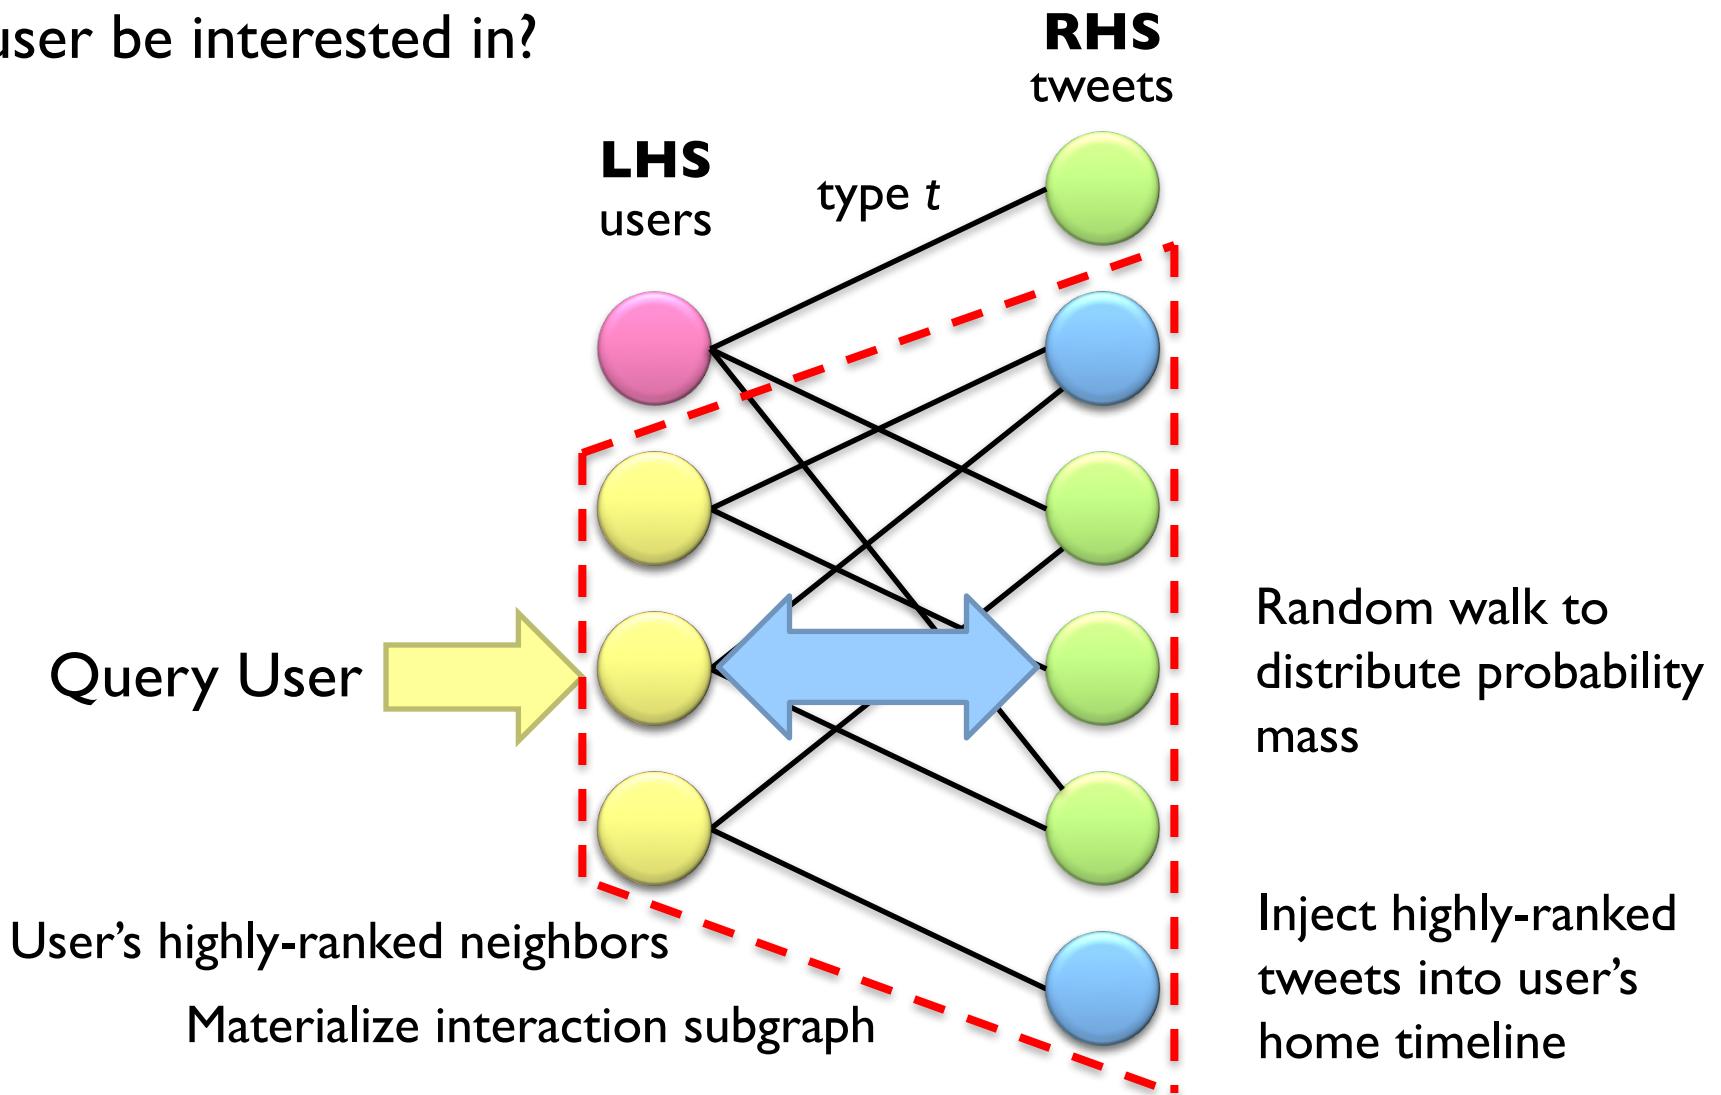
Index Segment

insertEdge


Moving Window

Recommendation Algorithm: Subgraph SALSA

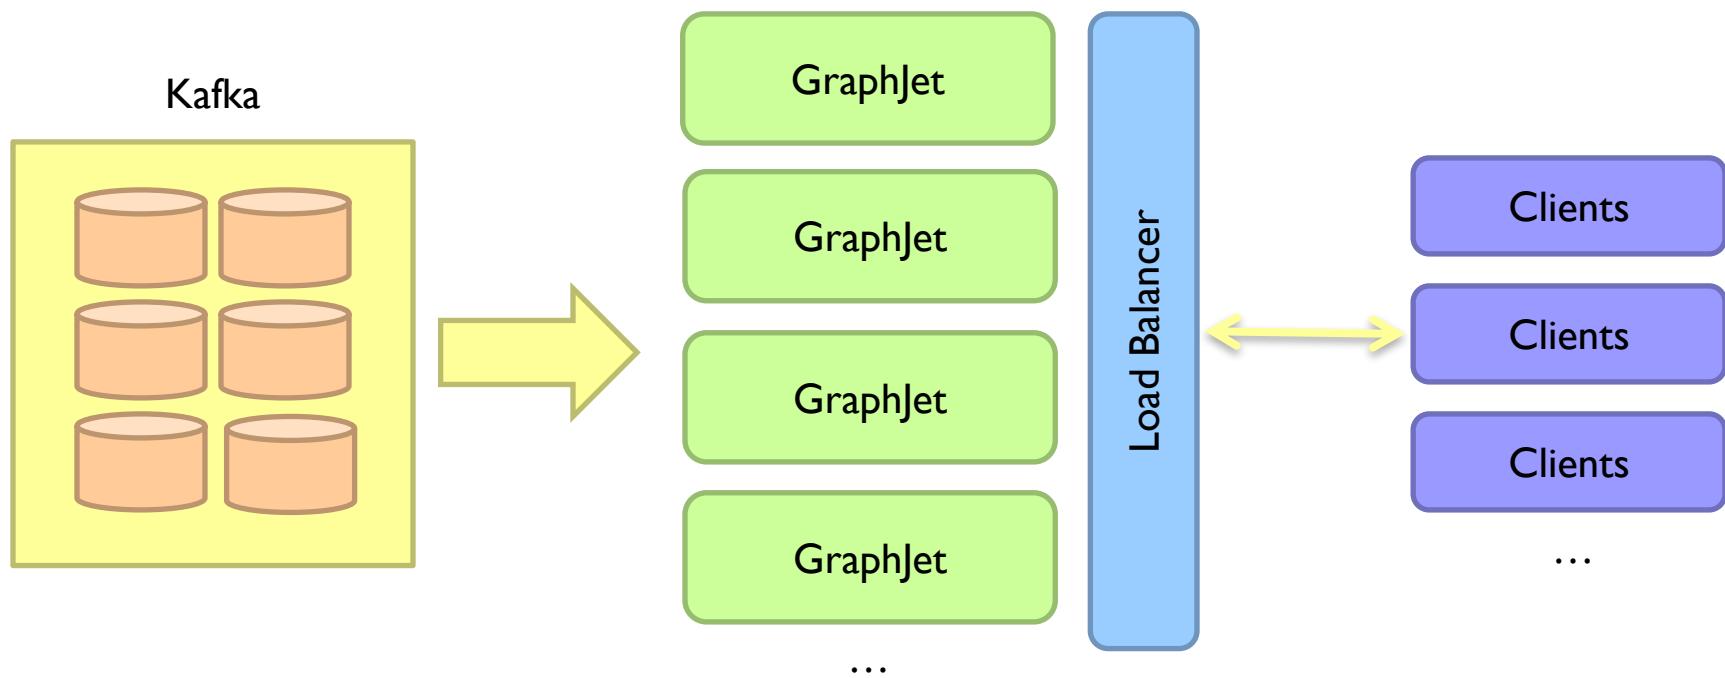
What tweets might a user be interested in?



Recommendation Algorithm: Similarity Query



Deployment Architecture



Production Status

Started serving production traffic early 2014

Dual Intel Xeon 6-cores (E5-2620 v2) at 2.1 GHz

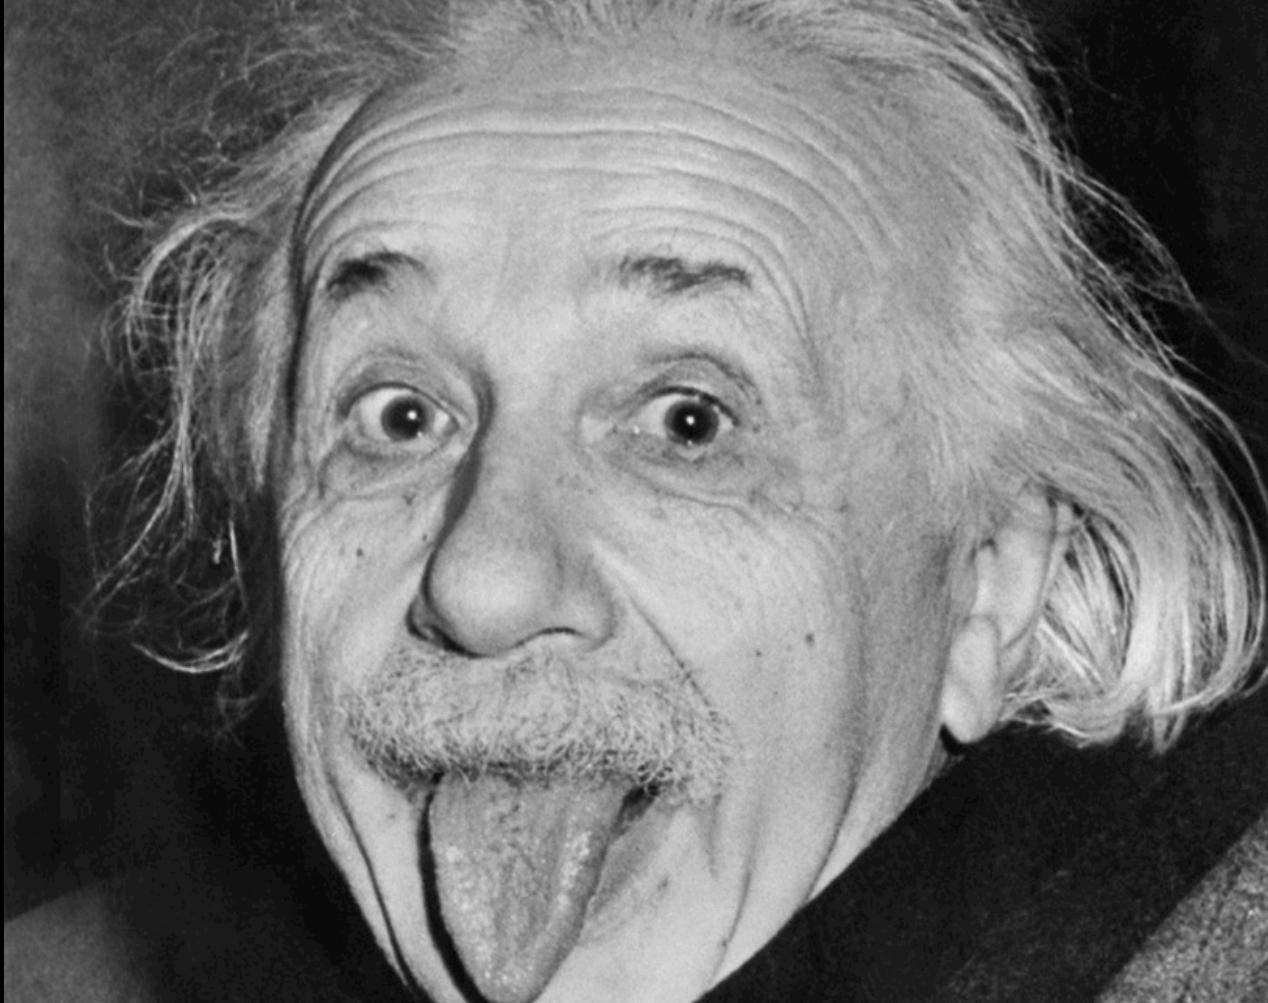
Cold startup: ingestion at $O(10^6)$ edges per sec from Kafka
Steady state: ingestion at $O(10^4)$ edges per sec

Space usage: $O(10^9)$ edges in < 30 GB

Sample recommendation algorithm: subgraph SALSA
500 QPS, p50 = 19ms, p99 = 33ms

Takeaway lesson #01:

Make things as simple as possible, but not simpler.



With lots of data, algorithms don't really matter that much
Why a complex architecture when a simple one suffices?



Takeaway lesson #10:
Constraints aren't always technical.



Takeaway lesson #11:

Visiting and revisiting design decisions

Questions?

“In theory, there is no difference between theory and practice. But, in practice, there is.”

- Jan L.A. van de Snepscheut

